


Product Demand Prediction with Machine Learnings

TEAM MEMBER

911721104089 : SEENU.A

Phase-1 Documentation submission

Project: Product Demand Prediction with Machine Learnings



OBJECTIVES:

Uninterrupted supply of products/services. Sales target setting and evaluating sales performance. Optimization of prices according to market fluctuations and inflation.

Phase 1: Problem Definition and Design Thinking

1.Data Collection:

Collect historical sales data and external factors that influence demand, such as marketing campaigns, holidays, economic indicators, etc.

Datalink:(<https://www.kaggle.com/datasets/chakradharmattapalli/product-demand-prediction-with-machine-learning>)

2.Data preprocessing

Data preprocessing is a crucial step in building machine learning models for product demand prediction. It involves cleaning, transforming, and organizing your data to make it suitable for training and testing your model. Here are the key steps in data preprocessing for product demand prediction:

Data Collection:

Gather historical data on product demand. This data should include information such as product attributes, time/date of sales, quantity sold, and any other relevant features.

Data Cleaning:

Handle missing data: Identify and decide how to handle missing values in your dataset. You can choose to remove rows with missing values, impute them with a suitable strategy (e.g., mean, median, or more advanced imputation methods), or treat them as a separate category.

Remove duplicates: Check for and remove any duplicate records in your dataset.

Feature Engineering:

Create new features: Generate additional features that might be relevant for demand prediction. For example, you can extract time-based features like day of the week, month, or season.

Encode categorical variables: Convert categorical variables (e.g., product categories) into numerical format using techniques like one-hot encoding or label encoding.

Data Transformation:

Scale features: Standardize or normalize numerical features to have a consistent scale, which can help some machine learning algorithms perform better.

Log transformations: If your target variable (demand) is highly skewed, applying a log transformation can make it more normally distributed, which is often beneficial for regression models.

Time Series Data:

If your data involves time series, consider handling it appropriately. This may include resampling to a consistent time interval, creating lag features, and dealing with seasonality and trends.

Train-Test Split:

Split your data into training and testing sets. The training set is used to train your machine learning model, while the testing set is used to evaluate its performance.

Outlier Detection and Handling:

Identify and handle outliers in your data. Outliers can have a significant impact on the performance of predictive models. You can choose to remove them or transform them using robust methods.

Feature Selection (Optional):

If you have many features, you may want to perform feature selection to choose the most relevant ones. Techniques like feature importance from tree-based models or dimensionality reduction methods can help.

Data Scaling and Normalization (Optional):

Depending on the machine learning algorithm you plan to use, scaling and normalization of features may be necessary. Some algorithms, like SVM and k-Nearest Neighbors, are sensitive to feature scales.

Data Pipeline:

Set up a data preprocessing pipeline to automate these steps. This ensures that the same preprocessing is applied to both the training and testing data consistently.

Validation and Cross-Validation:

Use cross-validation techniques to assess the performance of your model and prevent overfitting. This involves splitting your data into multiple folds for training and testing.

Save Processed Data:

After preprocessing, save your cleaned and transformed data to avoid having to repeat these steps every time you want to train or evaluate your model.

3.FUTURE ENGINEER'S:

Handle missing data: Identify and decide how to handle missing values in your dataset. You can choose to remove rows with missing values, impute them with a suitable strategy (e.g., mean, median, or more advanced imputation methods), or treat them as a separate category.

Remove duplicates: Check for and remove any duplicate records in your dataset.

Create new features: Generate additional features that might be relevant for demand prediction. For example, you can extract time-based features like day of the week, month, or season.

Encode categorical variables: Convert categorical variables (e.g., product categories) into numerical format using techniques like one-hot encoding or label encoding.

4.MODEL SELECTION:

- 1.decisiontree
- 2.linear regression

5.MODEL TRAINING:

Code :

```
import pandas as pd
import numpy as np
import plotly.express as px
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeRegressor

df=pd.read_csv('PoductDemand.csv')
```

6.EVALUATION:

Code:

```
import pandas as pd
import numpy as np
import plotly.express as px
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeRegressor

df=pd.read_csv('PoductDemand.csv')
print(df.head())
print(df.info())
print(df.describe())
print(df.isnull().sum())
df1=df.dropna()
print(df1)
print(df.corr())
x=df[["Total Price", "Base Price"]]
y=df["Units Sold"]
x_train, x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=42)
model=DecisionTreeRegressor(random_state = 0)
print(model.fit(x_train,y_train))
feature=np.array([[133.00,140.00]])
print(model.predict(feature))
plt.plot(x,y)
plt.xlabel('x')
plt.ylabel('y')
plt.show()
```

output:

```
ID Store ID Total Price Base Price Units Sold
0 1 8091 99.0375 111.8625 20
1 2 8091 99.0375 99.0375 28
2 3 8091 133.9500 133.9500 19
3 4 8091 133.9500 133.9500 44
4 5 8091 141.0750 141.0750 52
```

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 150150 entries, 0 to 150149

Data columns (total 5 columns):

```
# Column Non-Null Count Dtype
```

```
--- ----
```

```
0 ID 150150 non-null int64
```

```
1 Store ID 150150 non-null int64
2 Total Price 150149 non-null float64
3 Base Price 150150 non-null float64
4 Units Sold 150150 non-null int64
```

```
dtypes: float64(2), int64(3)
```

```
memory usage: 5.7 MB
```

```
None
```

```
      ID  Store ID  Total Price  Base Price  Units Sold
count 150150.000000 150150.000000 150149.000000 150150.000000 150150.000000
mean 106271.555504  9199.422511   206.626751   219.425927   51.674206
std  61386.037861  615.591445   103.308516   110.961712    60.207904
min    1.000000   8023.000000    41.325000    61.275000    1.000000
25%   53111.250000   8562.000000   130.387500   133.237500   20.000000
50%   106226.500000   9371.000000   198.075000   205.912500   35.000000
75%   159452.750000   9731.000000   233.700000   234.412500   62.000000
max   212644.000000   9984.000000   562.162500   562.162500  2876.000000
```

```
ID      0
```

```
Store ID      0
```

```
Total Price    1
```

```
Base Price     0
```

```
Units Sold     0
```

```
dtype: int64
```

```
      ID  Store ID  Total Price  Base Price  Units Sold
0      1    8091    99.0375   111.8625        20
1      2    8091    99.0375    99.0375        28
2      3    8091   133.9500   133.9500        19
3      4    8091   133.9500   133.9500        44
4      5    8091   141.0750   141.0750        52
...    ...    ...    ...    ...    ...
150145 212638    9984   235.8375   235.8375        38
150146 212639    9984   235.8375   235.8375        30
150147 212642    9984   357.6750   483.7875        31
150148 212643    9984   141.7875   191.6625        12
150149 212644    9984   234.4125   234.4125        15
```

```
[150149 rows x 5 columns]
```

```
      ID  Store ID  Total Price  Base Price  Units Sold
ID      1.000000  0.007464   0.008473   0.018932  -0.010616
Store ID  0.007464  1.000000  -0.038315  -0.038848  -0.004372
Total Price  0.008473 -0.038315   1.000000   0.958885  -0.235625
Base Price   0.018932 -0.038848   0.958885   1.000000  -0.140032
Units Sold  -0.010616 -0.004372  -0.235625  -0.140032   1.000000
DecisionTreeRegressor(random_state=0)
```

Linear regression

Figure 1

